

MVP Requirements

UI/FrontEnd:

1. Web App - **shadcn**
 1. Inputs:
 1. Model
 2. Config:
 1. Temperature (from 0 to 2)
 2. top-k (from _ to _)
 3. Context Window (tokens)
 4. Costs (\$)
 5. Chunk Size (input number)
 6. Tools (on/off)
 7. Max tokens
 3. the same prompts (temp) - a fixed list of 200 prompts (all kinds of prompts = long/short, with/without tool use)
 2. Output:
 1. dashboard - failure rate, failure confidence band, description

BackEnd:

1. variables are accounted

Task division:

1. Person A: Chris
2. Person B: Emil
3. Person C: Yufei

Probabilistic Failure Simulator Hackathon Work Plans

Brief overview: Three coordinated tracks build a simulator that runs probes, logs failures, computes probabilities, and visualizes results. Plans include roles, responsibilities, timelines, dependencies, files, success criteria, and handoffs.

Person A: Runner + Rules Engineer

Role summary: Build probe runner, telemetry, and deterministic rules to generate failure events and break-first timeline.

Core responsibilities:

- Implement probe execution against configs.
- Log tokens, latency, tool calls.
- Apply rules to emit failure events and break-first timeline.
 - Key files/folders:
 - /lib/probe-runner.ts
 - /lib/rules-engine.ts
 - /lib/telemetry-logger.ts
 - Function signatures:
 - runProbe(config: ProbeConfig): ProbeResult
 - loadConfigs(dir: string): ProbeConfig[]
 - logTelemetry(t: TelemetryRecord): void
 - evaluateRules(r: ProbeResult, rules: Rule[]): FailureEvent[]
 - buildBreakFirstTimeline(events: FailureEvent[]): Timeline
 - Dependencies:
 - Inputs: /configs/*.json
 - Outputs consumed by B and C.
 - Handoff points:
 - failure-events.json, break-first-timeline.json, telemetry.log
 - Success criteria:
 - Deterministic, reproducible events.
 - Telemetry completeness (>98% probes logged).
 - Break-first timeline aligns with events.
 - Saturday timeline (hour-by-hour):
 - H0–1: Repo setup, shared schema sync (all).
 - H2: Scaffolding files; ProbeConfig schema (/types.ts).
 - H3: loadConfigs(), runProbe() stub; sample configs.
 - H4: logTelemetry() with token/latency/tool fields; write telemetry.log.
 - H5: Implement evaluateRules() with rule DSL.
 - H6: Emit failure-events.json {probeId, configId, mode, ts}.
 - H7: buildBreakFirstTimeline(); write break-first-timeline.json.
 - H8: CLI: npm run probes -> outputs all files.
 - H9: Add seed for determinism; doc handoff format.

- H10–12: Milestone: freeze v1 outputs and schemas; push.
Sunday timeline:
- H1: Bugfix runner edge cases; add retries.
- H2: Add more failure modes; rule examples.
- H3: Perf pass; batch I/O.
- H4: Schema lock; write README section for A.
- H5: Validate with B; sample datasets.
- H6: Feature freeze (all).
- H7–8: Final submission packaging.

Person B: Probability + Analytics Engineer

Role summary: Compute per-config failure probabilities, uncertainty, safer comparisons, and distributions.

Core responsibilities:

- \hat{p} per config; 95% CI (Bootstrap/Bayesian).
 - $P(A \text{ safer than } B)$.
 - Distributions by failure mode/prompt family.
- Key files/folders:
- /lib/probability.ts
 - /lib/statistics.ts
 - /lib/analysis.ts
- Function signatures:
- estimatePhat(events: FailureEvent[], configId: string): {k: number, n: number, phat: number}
 - bootstrapCI(k: number, n: number, alpha = 0.05): [number, number]
 - bayesianBetaCI(k: number, n: number, alpha = 0.05): [number, number]
 - compareConfigs(a: Stats, b: Stats): {pASafer: number}
 - modeDistributions(events: FailureEvent[]): Record
- Dependencies:
- Consumes: failure-events.json, telemetry.log.
 - Provides to C.
- Handoff points:
- analysis.json, distributions.json, comparisons.json
- Success criteria:
- Correct CIs; unit tests for edge cases.
 - Clear safer-than metrics.
 - Fast (<1s per dataset).
- Saturday timeline:
- H0–1: Schema sync (all).
 - H2: Parse events; define Stats type.
 - H3: estimatePhat(); write analysis.json (per config).
 - H4: bootstrapCI(), bayesianBetaCI().
 - H5: compareConfigs(); write comparisons.json.
 - H6: modeDistributions(); write distributions.json.
 - H7: Family aggregation; prompt-family.json.
 - H8: Tests (/tests/statistics.test.ts).
 - H9: Perf/validation against synthetic data.

- H10–12: Milestone: freeze API and outputs.
Sunday timeline:
- H1: Edge cases (zero events); fallbacks.
- H2: Add CLI: npm run analyze.
- H3: Docs of JSON schemas.
- H4: Align with C on visualization shapes.
- H5: Final QA with A datasets.
- H6: Feature freeze (all).
- H7–8: Submission artifacts.

Person C: Frontend + Story Engineer

Role summary: Build UI to input configs, show analytics, timelines, breakdowns; export; demo and README.

Core responsibilities:

- Scaffold app with v0, shadcn/ui.

- Cards for Config A/B.

- Dashboard for distributions and timelines.

- Export JSON/Markdown; optional n8n.

Key files/folders:

- /app

- /components

- README.md

- /public/demo

Components and functions:

- ConfigCard.tsx

- Dashboard.tsx

- TimelineChart.tsx

- DistributionChart.tsx

- loadData(files: string[]): AppData

- exportJSON(data: AppData): File

- exportMarkdown(summary: Summary): File

Dependencies:

- Consumes A/B JSONs.

- Requires stable schemas from B.

Handoff points:

- UI consumes analysis.json, comparisons.json, distributions.json, break-first-timeline.json

Success criteria:

- Clear comparisons and timelines.

- Export works.

- Demo video shows end-to-end.

Saturday timeline:

- H0–1: Design and data schema sync (all).

- H2: Next.js + shadcn scaffold (/app).

- H3: ConfigCard.tsx; file picker.

- H4: Data loader; schema validation.

- H5: Dashboard layout; cards/sections.

- H6: TimelineChart wired to break-first-timeline.json.

- H7: DistributionChart from distributions.json.
- H8: Safer-than widget using comparisons.json.
- H9: Export JSON/MD.
- H10–12: Milestone: UI v1 walkthrough.
Sunday timeline:
- H1: Polish UI; accessibility.
- H2: n8n integration stub (optional).
- H3: README.md authoring; usage.
- H4: Record demo video.
- H5: Bugfixes; perf.
- H6: Feature freeze (all).
- H7–8: Final submission packaging.